Weak Scaling on Leonardo

In **weak scaling**, both the **number of computing units** and the **problem size** are **increased**: *constant workload per computing unit*.

We use 8×10^6 unknowns per GPU, *i.e.*, 3.2×10^7 unknowns per node.

We use the following resources:

- Number of GPUs from 1 to 8192,
- GPUs x Node 4 (1 MPI Task x GPU, 8 CPUs per Task)
- Pure MPI: 32 MPI Tasks per Node

Within the software framework:

- Compilers: `gcc/11.3.0`
- MPI: `openmpi/4.1.4`
- CUDA `compilation tools, release 11.8, V11.8.89`